

Process Management

Process Management

- Processes
 - Threads
 - Interprocess Communication (IPC)
 - CPU Scheduling
 - Process Synchronization
 - Deadlocks

Processes

A *process* is a set of identifiable, repeatable actions which are ordered in some way and contribute to the fulfillment of an objective.

(General definition)

A *process* is a program in execution.

(Computer oriented definition)

- **Program: static, passive**

A cooking recipe is a program.

- **Process: dynamic, active**

Acting according to the recipe (cooking) is a process.

BOBBY DARIN'S MANICOTTI
 1 package Juliette Manicotti
 3 lbs fresh ricotta
 2 packages Mozzarella
 3 eggs
 2 lbs chopped beef
 8 cloves garlic
 2 onions
 1/2 teaspoon salt
 10 cans Hunts or Del Monte tomato sauce
 1 teaspoon oregano
 1 teaspoon basil
 1 cup brea
 1/2 lb grated
 1/2 cup oil
 1/2 teaspoon
 Make sau
 onions and g
 oregano, basi




Process Model

Processes

- Several processes are working *quasi*-parallel.
- A *process* is a unit of work.
- Conceptually, each process has its own virtual CPU.


In reality, the real CPU switches back and forth from process to process.

One program counter



Sequential view

Four program counters



Process model view




Figure from [Ta01 p.72]

Processes

A process may be described as either (more or less)

a) CPU-bound

spends more time doing computations – few very long CPU bursts.

b) I/O-bound

spends more time doing I/O than computations – many short CPU bursts.




Figure from [Ta01 p.134]

Address Space

Processes

A *process* is an executing program, and encompasses the current values of the program counter, of the registers, of the variables and of the stack.

- **code section (text section or *segment*)**


This is the actual program code (the machine instructions).

- **data section (data segment)**

This segment contains global variables (global to the process, not global to the computer system).

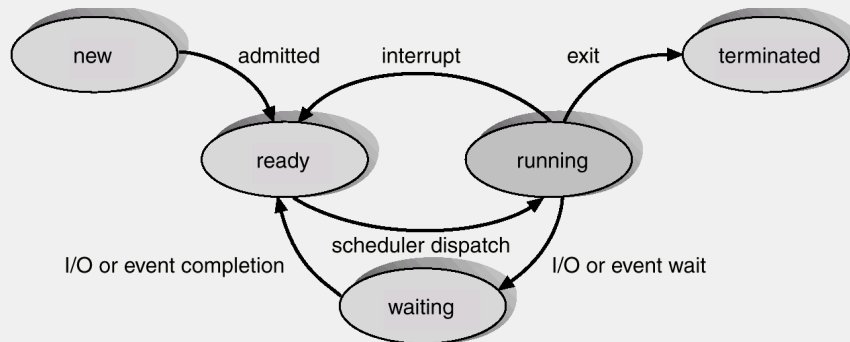
- **stack section (stack segment)**

The stack contains temporary data (local variables, return addresses, function parameters).



Process States

Processes



Note: Only in the *running* state the process needs CPU cycles, in all other states it is actually 'frozen' (or nonexistent any more).

Figure from [Sil00 p.89]

Process States

Processes

- **New**

The process is created. Resources are allocated.

- **Ready**

The process is ready to be (re)scheduled.

- **Running**


The CPU is allocated to the process, that is, the program instructions are being executed.

- **Waiting**

The process is waiting for some event to occur. Without this event the process cannot continue – even if the CPU would be free.

- **Terminated**

Work is done. The process is taken off the system (off the queues) and its resources are freed.



Processes

Events at which processes are created

- **Operating System Start-Up**

Most of the system processes are created here. A large portion of them are background processes (daemons).

- **Interactive User Request**

A user requests an application to start.

- **Batch job**

Jobs that are scheduled to be carried out when the system has available the resources (e.g. calendar-driven events, low priority jobs)

- **Existing process gives birth**

An existing process (e.g. a user application or a system process) creates a process to carry out some related (sub)tasks.

Process Creation

Processes

- **Parent process creates a child process**

which in turn may create other processes, forming a tree of processes.

- **Resource sharing**

- Parent and child share all resources.
- Child shares subset of parent's resources.
- Parent and child share no resources.

- **Execution**

- Parent and child execute concurrently.
- Parent waits until child terminates.

- **Address Space**

- Child is copy of parent
- Child has program loaded into it

System calls for creating a child process:
Unix: `fork()`
Windows: `CreateProcess()`

fork () example

Process Creation

```
#include <stdio.h>

void main()
{
    int result;
    printf("Parent, my pid is: %d\n", getpid());
    result = fork();
    if (result == 0) { /* child only */
        printf("Child, my pid is: %d\n", getpid());
        ...
    } else { /* parent only */
        printf("Parent, my pid is: %d\n", getpid());
        ...
    }
}
```

system call that tells a process its `pid` (process identifier) which is a unique process number within the system.

Executed by child

Executed by parent

from here on think parallel

fork () example

Process Creation


Terminal output: Parent, my pid is: 189

Child, my pid is: 190



Parent, my pid is: 189

order depends on whether parent or child is scheduled first after `fork()`.

Before `fork()`



After `fork()`



Process Creation

Processes

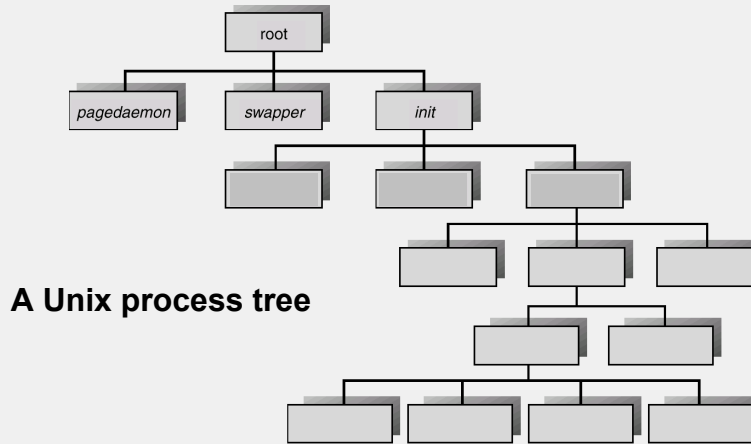


Figure from [Sil00 p.96]

Process Termination

Processes

Events at which processes are terminated

- Process asks the OS to delete it

Work is done. Resources are deallocated (memory is freed, open files are closed, used I/O buffers are flushed).

- Parent terminates child

- Child may have exceeded allocated resources.
- Task assigned to child is no longer required.
- Parent's parent is exiting.

Some OS do not allow a child to continue when its parent terminates. Cascading termination (a sub tree is deleted).

System calls for self-termination of a process:
Unix: `exit()`
Windows: `ExitProcess()`

Process Control Block

Processes

- Operating system maintains a process table
- Each entry represents a process
- Entry often termed *PCB* (process control block)

A PCB contains all information about a process that must be saved when the process is switched from *running* into *waiting* or *ready*, such that it can later be restarted as if it had never been stopped.

- Info regarding process management,
- regarding memory occupation
- and open files.

Figure from [Sil00 p.89]

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

PCB example

Process Control Block

Processes

Process management

Registers
Program counter
Program status word
Stack pointer
Process state
Priority
Scheduling parameters
Process ID
Parent process
Process group
Signals
Time when process started
CPU time used
Children's CPU time
Time of next alarm

Memory management

Pointer to text segment
Pointer to data segment
Pointer to stack segment

File management

Root directory
Working directory
File descriptors
User ID
Group ID

Typical fields of a PCB

Table from [Ta01 p.80]

Context Switch

Processes

The task of switching the CPU from one process to another is termed *context switch* (sometimes also *process switch*):

- Saving the state of old process

Saving the current *context* of the process in its PCB.

- Loading the state of new process

Restoring the former context of the process from its PCB.

Context switching is pure administrative overhead. The duration of a switch lies in the range of 1 ... 1000 μ s. The switch time depends on the hardware. Processors with multiple sets of registers are faster in switching. Context switching poses a certain bottleneck, which is one reason for the introduction of *threads*.

Context Switch

Processes

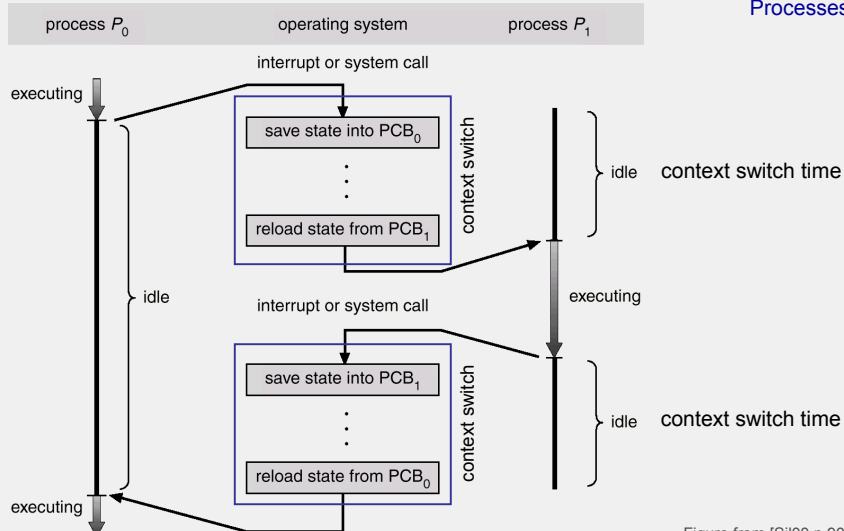


Figure from [Sil00 p.90]

Scheduling

Processes

On a uniprocessor system there is only one process running, all others have to wait until they are scheduled. They are waiting in some **scheduling queue**:

- **Job Queue**

Holds the future processes of the system.

- **Ready Queue** (also called **CPU queue**)

Holds all processes that reside in memory and are ready to execute.

- **Device Queue** (also called **I/O queue**)

Each device has a queue holding the processes waiting for I/O completion.

- **IPC Queue**

Holds the processes that wait for some IPC (inter process communication) event to occur.