

Mutual Exclusion

IPC

Intermediate Summary

- | | |
|------------------------|--|
| • Disabling Interrupts | Not recommended for multi-user systems. |
| • Lock Variables | Problem remains the same. |
| • Strict Alternation | Violation of condition 3. Busy waiting. |
| • Peterson Algorithm | Busy waiting. |
| • TSL instruction | Solves the problem through atomic operation. Should be used without busy waiting. |

In essence, what the last three solutions do is this: A process checks whether the entry to its critical region is allowed. If it is not, the process just sits in a tight loop waiting until it is.

Unexpected side effects, such as *priority inversion problem*.

Priority Inversion Problem

IPC

Consider a computer with two processes

Process **H** with high priority

Process **L** with low priority

The scheduling rules are such that **H** is run whenever it is in ready state.

At a certain moment, with **L** in its critical region, **H** becomes ready and is scheduled. **H** now begins busy waiting, but since **L** is never scheduled while **H** is running, **L** never has the chance to leave its critical region. **H** loops forever. This is sometimes referred to as the *priority inversion problem*.

Solution: *blocking* a process instead of wasting CPU time.

Sleep and wake up

IPC

sleep()

A system call that causes the caller to block, that is, the process voluntarily goes from the *running* state into the *waiting* state. The scheduler switches over to another process.

wakeup(process)

A system call that causes the process *process* to awake from its **sleep()** and to continue execution. If the process *process* is not **asleep** at that moment, the wakeup signal is lost.

Note: these two calls are fictitious representatives of real system calls whose names and parameters depend on the particular operating system.

Producer – Consumer Problem

IPC

- **Shared buffer with limited size**


The buffer allows for a maximum of N entries (it is bounded). The problem is also known as *bounded buffer* problem.

- **Producer puts information into buffer**

When the buffer is full, the producer must wait until at least one item has been consumed.

- **Consumer removes information from buffer**

When the buffer is empty the consumer must wait until at least one new item has been entered.



```
const int N = 100;
int count = 0;
```

Producer – Consumer Implementation Example
This implementation suffers from race conditions

```
void producer() {
    while (TRUE) {
        int item = produce_item();
        if (count == N) sleep();
        insert_item(item);
        count++;
        if (count == 1) wakeup(consumer);
    }
}

void consumer() {
    while (TRUE) {
        if (count == 0) sleep();
        item = remove_item();
        count--;
        if (count == N-1) wakeup(producer);
        consume_item(item);
    }
}
```

Annotations for Producer:

- constantly producing
- produce item
- sleep when buffer is full
- enter item to buffer
- adjust item counter
- wakeup(consumer);

Annotation for Consumer:

- constantly consuming
- sleep when buffer is empty
- remove one item
- adjust item counter
- wakeup(producer);
- consume_item(item);

Diagram A:

A circle labeled 'A' is positioned between the two functions. A line connects it to the `if (count == 0) sleep();` line in the consumer function. Another line connects it to the `if (count == 1) wakeup(consumer);` line in the producer function.

Annotations for Diagram A:

- when the buffer was empty beforehand (and thus now has 1 item), wakeup any consumer(s) that may be waiting
- when the buffer was full beforehand (and thus now has N-1 items), wakeup producer(s) that may be waiting.

Producer – Consumer Problem

A race condition may occur in this case:

Mutual Exclusion

The buffer is empty and the consumer has just read `count` to see if it is 0. At that instant (see (A) in listing) the scheduler decides to switch over to the producer.

The producer inserts an item in the buffer, increments `count` and notices that `count` is now 1. Reasoning that `count` was just 0 and thus the consumer must be sleeping, the producer calls `wakeup()` to wake the consumer up.

However, the consumer was not yet asleep, it was just taken away the CPU shortly before it could enter `sleep()`. The wakeup signal is lost.

When the consumer is rescheduled and resumes at (A), it will go to sleep. Sooner or later the producer has filled up the buffer and goes asleep as well.

Both processes will sleep forever.

Producer – Consumer Problem

Mutual Exclusion

Reasons for race condition

- The variable `count` is unconstrained
Any process has access any time.
- Evaluating `count` and going asleep is a non-atomic operation
The prerequisite(s) that lead to `sleep()` may have changed when `sleep()` is reached.

Workaround:

- Add a wakeup waiting bit
When the bit is set, `sleep()` will reset that bit and the process stays awake.
- Each process must have a wakeup bit assigned
Although this is possible, the principal problem is not solved.

What is needed is something that does *testing a variable and going to sleep* – dependent on that variable – in a single non-interruptible manner.

Semaphores

Mutual Exclusion

- Introduced by Dijkstra (1965)
- Counting the number of wakeups

An integer variable counts the number of wakeups for future use.

- Two operations: **down** and **up**

down is a generalization of **sleep**. **up** is a generalization of **wakeup**. Both operations are carried out in a **single, indivisible operation** (usually in kernel). Once a semaphore operation is started, no other process can access the semaphore.

```
down(int* sem) {  
    if (*sem < 1) sleep();  
    *sem--;  
}
```

principle of **down**-operation

```
up(int* sem) {  
    *sem++;  
    if (*sem == 1) wakeup a process  
}
```

principle of **up**-operation

Semaphores

Mutual Exclusion

- **Up and down are system calls**

in order to make sure that the operating system briefly disables all interrupts while carrying out the few machine instructions implementing **up** and **down**.

- **Semaphores should be lock-protected**

This is recommended at least in multi-processor systems to prevent another CPU from simultaneously accessing a semaphore. TSL instruction helps out.

Producer – Consumer problem using semaphores (next page)

Definition of variables:

| | |
|-----------------------------------|-------------------------------------|
| | <code>const int N = 10;</code> |
| a semaphore is an integer | <code>typedef int semaphore;</code> |
| counting empty slots | <code>semaphore empty = N;</code> |
| counting full slots | <code>semaphore full = 0;</code> |
| mutual exclusion on buffer access | <code>semaphore mutex = 1;</code> |

Computer Architecture WS 06/07 Dr.-Ing. Stefan Freinatis



Producer – Consumer Implementation Example
This implementation does not suffer from race conditions


```
void producer() {
    while (TRUE) {
        int item = produce_item();
        down(&empty);           possibly sleep, decrement empty counter
        down(&mutex);           possibly sleep, claim mutex (set it to 0) thereafter
        insert_item(item);
        up(&mutex);             release mutex, wake up other process
        up(&full);              increment full counter, possibly wake up other ...
    }
}

void consumer() {
    while (TRUE) {
        down(&full);           possibly sleep, decrement full counter
        down(&mutex);           possibly sleep, claim mutex (set it to 0) thereafter
        item = remove_item();
        up(&mutex);             release mutex, wake up other process
        up(&empty);            increment empty counter, possibly wake up other ...
        consume_item(item);
    }
}
```

Semaphores


Assume $N = 5$. Initial condition: **empty** = 5, **full** = 0.

Mutual Exclusion



Scenario: producer is working, no consumer present

Initial condition: **empty** = 0, **full** = 5.




Scenario: consumer is working, no producer present

Semaphores

Assume $N = 5$. Initial condition: **empty** = 1, **full** = 4.

Mutual Exclusion




Scenario: Consumer waking up producer

Semaphores

Assume $N = 5$. Initial condition: **empty** = 4, **full** = 1.

Mutual Exclusion

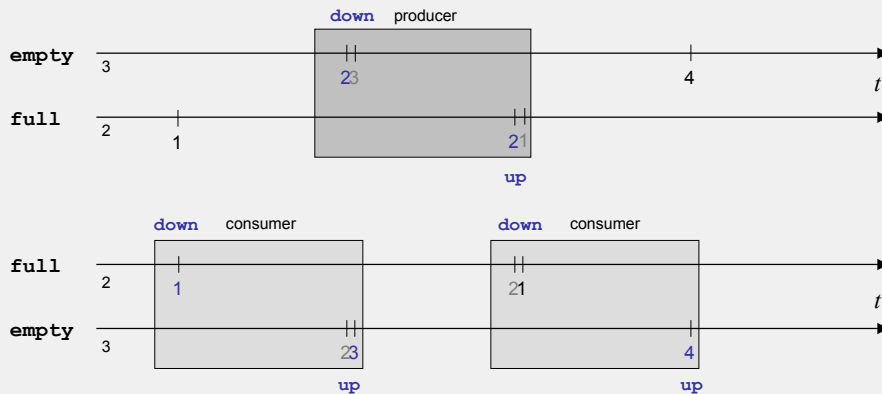


Scenario: Producer waking up consumer

Semaphores

Assume $N = 5$. Initial condition: **empty** = 3, **full** = 2.

Mutual Exclusion



If processes overlap, then temporary it may be that **empty** + **full** \neq N

Note that consumer and producer may almost concurrently change the same semaphore legally.

Mutex

Mutual Exclusion

- **Simplified semaphore**

when counting is not needed.

- **Two states**

Locked or unlocked. Used for managing mutual exclusion (hence the name).

| | | | |
|----------------------------|--------------------------------|---|---------------------------|
| <code>mutex_lock:</code> | <code>TSL R, mutex</code> | — | get and set mutex |
| | <code>CMP R, #0</code> | — | was it unlocked? |
| | <code>JZ ok</code> | — | if yes: jump to ok |
| | <code>CALL thread_yield</code> | — | if no: sleep |
| | <code>JMP mutex_lock</code> | — | try again acquiring mutex |
| <code>ok:</code> | <code>RET</code> | | |
| <code>mutex_unlock:</code> | <code>MOV mutex, #0</code> | — | unlock mutex |
| | <code>RET</code> | | |

Pseudo assembler listing implementing `mutex_lock()` and `mutex_unlock()`.

Monitors

Mutual Exclusion

- **High level synchronization primitive**

at programming language level. Direct support by some programming languages.

- **A collection of procedures, variables and data structures grouped together in a module**

- A monitor has multiple entry points
- Only one process can be in the monitor at a time
- Enforces mutual exclusion – less chances for programming errors

- **Monitor implementation**

- Compiler handles implementation
- Library functions using semaphores

Monitors

Mutual Exclusion

```
monitor example;  
  integer i;  
  condition c;
```

Variables not accessible from outside the monitor's own methods (capsulation).

```
  procedure producer()  
  ...  
  end;
```

Functions (methods) publicly accessible to all processes, however only one process at a time may call a monitor function.

```
  procedure consumer()  
  ...  
  end;
```

If the buffer is full, the producer must wait.
If the buffer is empty the consumer must wait.

```
end monitor;
```

A monitor in Pidgin Pascal, from [Ta01 p.115]

Monitors

Mutual Exclusion

- How can a process wait inside a monitor?

Cannot put to sleep because no other process can enter the monitor meanwhile.

- Use a condition variable!

A condition variable supports two operations.

- **wait()**: suspend this process until it is signaled. The suspended process is not considered *inside* the monitor any more. Another process is allowed to enter the monitor.
- **signal()**: wake up one process waiting on the condition variable. No effect if nobody is waiting. The signaling process automatically leaves the monitor (*Hoare* monitor).
- Condition variables usable only inside a monitor.

Monitors

Mutual Exclusion

```
monitor ProducerConsumer
condition full, empty;
integer count;
procedure insert(item: integer);
begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
end;
function remove: integer;
begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
end;
count := 0;
end monitor;
```

```
procedure producer;
begin
    while true do
    begin
        item = produce_item;
        ProducerConsumer.insert(item)
    end
end;
procedure consumer;
begin
    while true do
    begin
        item = ProducerConsumer.remove;
        consume_item(item)
    end
end;
```

Producer-Consumer problem with monitors, from [Ta01 p.117]

Barriers

IPC

- Group synchronization
 - Intended for groups of processes rather than for two processes.
- Processes wait at a *barrier* for the others
 - according to the *all-or-none* principle
- After all have arrived, all can proceed


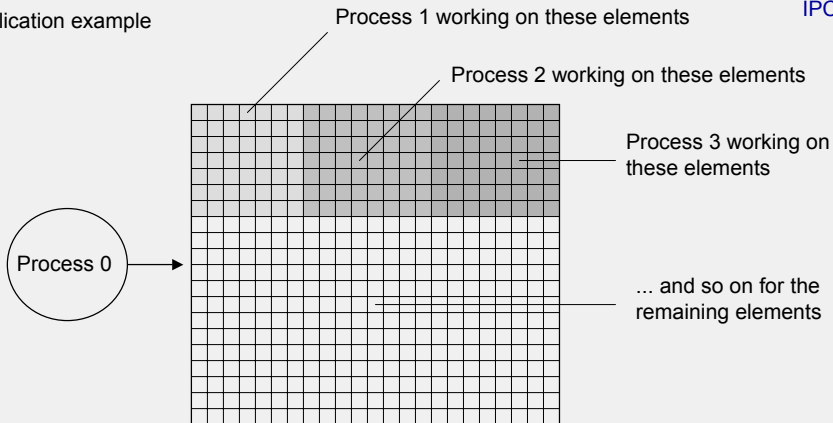


Figure from [Ta01 p.124]

Barriers

Application example

IPC



An array (e.g. an image) is updated frequently by some process 0 (producer). Many processes are working in parallel on certain array elements (consumers). All consumers must wait until the array has been updated and can then start working again on the updated input.

IPC

Intermediate Summary (II)

- **Semaphores**

Counting variable, used in non-interruptible manner. **Down** may put the caller to sleep, **up** may wake up another process.

- **Mutexes**

Simplified semaphore with two states. Used for mutual exclusion.

- **Monitors**

High level construct for achieving mutual exclusion at programming language level.

- **Barriers**

Used for synchronizing a group of processes.

These mechanisms all serve for process synchronization.

For data exchange among processes something else is needed: Messages.

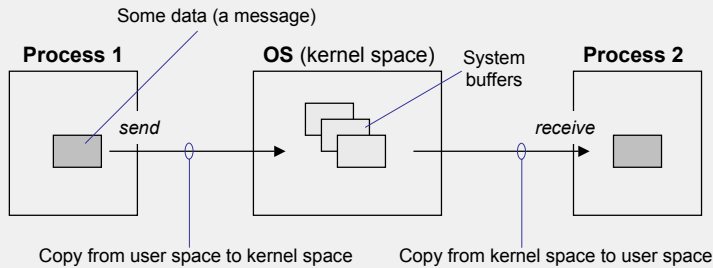
Messages

IPC

- Kernel supported mechanism for data exchange

Eliminates the need for 'self-made' (user-programmed) communication via shared resources such as shared files or shared memory.

- Two basic operations:
 - **send()**: send data
 - **receive()**: receive data



Direct Communication

Messages

- Both processes must exist

As the name *direct* implies, you cannot send a message to a future process.

- Processes must name each other explicitly

- `send(P, message)`: send data to process P
- `receive(Q, message)`: receive data from process Q

Symmetry in addressing. Both processes need to know each other by some identifier. This is no problem if both were `fork()`ed off the same parent beforehand, but is a problem when they are 'strangers' to each other.

- Communication link properties

- One process pair has exactly one link
- The link may be unidirectional or bidirectional

Indirect Communication

Messages

- Messages are send / received from mailboxes

The mailbox must exist, not necessarily the receiving process yet.

- Each mailbox has a unique identifier
- Processes communicate when they access the same mailbox

- Primitives

- send(A, message): send message to mailbox A
- receive(A, message): receive message from mailbox A

- Communication link properties


- Link is established when processes share a mailbox
- A link may be associated with many processes (broadcast)
- Unidirectional or bidirectional communication

Synchronous Communication

Messages

- Also called *blocking* send / receive
- Sender waits for receiver to receive the data

The send() system call blocks until receiver has received the message.



- Receiver waits for sender to send data


The receive() system call blocks until a message is arriving.

Asynchronous Communication

Messages

- Also called *non-blocking* send / receive
- Sender drops message and passes on

The send() system call returns to the caller when the kernel has the message.



- Receiver peeks for messages

The receive() system does not block, but rather returns some error code telling whether there is a message or not. Receiver must do *polling* to check for messages.

Messages

IPC

- Send by copy

The message is copied to kernel buffer at send time. At receive time the message is copied to the receiver. Copying takes time.

- Send by reference

A reference (a memory address or a handle) is copied to the receiver which uses the reference to access the data. The data usually resides in a kernel buffer (is copied there beforehand). Fast read access.

- Fixed sized messages

The kernel buffers are of fixed size – as are the messages. Straightforward system level implementation. Big messages must be constructed from many small messages which makes user level programming somewhat more difficult.

- Variable sized messages

Sender and receiver must communicate about the message size. Best use of kernel buffer space, however, buffers must not grow indefinitely.

UNIX IPC Mechanisms

IPC

- Pipes

Simple(st) communication link between two processes. Applies *first-in first-out* principle. Works like an invisible file, but is no file. Operations: `read()`, `write()`.

- FIFOs

Also called *named* pipe. Works like a file. May exist in modern Unices just in the kernel (and not in the file system). There can be more than one writer or reader on a FIFO. Operations: `open()`, `close()`, `read()`, `write()`.

- Messages

Allow for message transfer. Messages can have types. A process may read all messages or only those of a particular type. Message communication works according to the *first-in first-out* principle. Operations: `msgget()`, `msgsnd()`, `msgrcv()`, `msgctl()`.

UNIX IPC Mechanisms

IPC

- Shared memory

A selectable part of the address space of process P_1 is mapped into the address space of another process P_2 (or others). The processes have simultaneous access. Operations: `shmget()`, `shmat()`, `shmdt()`, `shmctl()`.

- Semaphores

Creation and manipulation of sets of semaphores. Operations: `semget()`, `semop()`, `semctl()`.

For an introduction into the UNIX IPC mechanisms (with examples) see

Stefan Freinatis: *Interprozeßkommunikation unter Unix - eine Einführung*,
Technischer Bericht, Fachgebiet Datenverarbeitung, Universität Duisburg, 1994.

<http://www.fb9dv.uni-duisburg.de/vs/members/fr/ipc.pdf>

Simple **pipe** example. Parent is writing, child is reading.

```
const int FIXSIZE=80
void main() {
    int fd[2];                // file descriptors for pipe
    pipe(fd);                 // create pipe
    int result = fork();      // duplicate process
    if (result == 0) {        // start child's code
        printf("This is the child, my pid is: %d\n", getpid());
        close(fd[1]);         // we do not need writing
        char buf[256];        // a buffer
        read(fd[0], buf, FIXSIZE) // wait for message from parent
        printf("Child: received message was: %s\n", buf);
        exit(0);              // good bye
    }                          // end child, start parent
    close(fd[0]);             // we do not need reading
    printf("This is the parent, my pid is: %d\n", getpid());
    write(fd[1], "Hallo!", FIXSIZE); // write message to child
}
```

Classical IPC Problems

The *dining philosophers*

IPC

An artificial synchronization problem posed and solved by Edsger Dijkstra 1965.

- Five philosophers sitting at a table

The problem can be generalized to more than five philosophers, of course.

- Each either eats or thinks
- Five forks available
- Eating needs 2 forks
- Pick one fork at a time

Slippery spaghetti, one needs two forks!

Either first the right fork and then the left one, or vice versa.




Figure from [Ta01 p.125]

Dining philosophers

Classical IPC problems

The life of these philosophers consists of alternate periods of eating and thinking. When a philosopher becomes hungry, she tries to acquire her left and right fork, one at a time, in either order. If successful in acquiring two forks, she eats for a while, then puts down the forks and continues to think.

Text from [Ta01 p.125]

Can you write a program that

- makes the philosophers eating and thinking (thus creation of 5 threads or processes, one for each philosopher),
- allows maximum utilization (parallelism), that is, two philosophers may eat at a time (no simple solution with just one philosopher eating at a time),
- is not centrally controlled by somebody instructing the philosophers,
- and that never gets stuck?

Dining philosophers

Classical IPC problems

```
const int N=5;

void philosopher(int i) {           // N philosophers in parallel
    while(TRUE) {                  // for the whole life
        think();
        take_fork(i);              // take left fork
        take_fork((i+1)%N);        // take right fork
        eat();
        put_fork(i);               // put left fork
        put_fork((i+1)%N);        // put right fork
    }
}
```

A *nonsolution* to the dining philosophers problem

If all philosophers take their left fork simultaneously, none will be able to take the right fork. All philosophers get stuck. Deadlock situation.

Sleeping Barber

IPC

The barber shop has one barber, one barber chair, and n chairs for customers, if any, to sit on. If there are no customers present, the barber sits down in the barber chair and falls asleep. When a customer arrives, he has to wake up the sleeping barber. If additional customers arrive while the barber is cutting a customer's hair, they either sit down (if there are empty chairs) or leave the shop (if all chairs are full).

Text from [Ta01 p.129]

How to program the barber and the customers without getting into race conditions?

```
const int CHAIRS=5;          // number of chairs
typedef int semaphore;
semaphore customers = 0;     // number of customers waiting
semaphore barbers = 0;      // number of barbers waiting
semaphore mutex = 1;        // for mutual exclusion
int waiting = 0;
```

```
void barber() {              // barber process
    while(TRUE){             // for the whole life
        down(&customers);     // sleep if no customers
        down(&mutex);         // acquire access to 'waiting'
        waiting--;
        up(&barbers);         // one barber ready to cut
        up(&mutex);           // release 'waiting'
        cut_hair();           // cut hair (non critical)
    }
}
```

A solution to the sleeping barber problem [Ta01 p.131]

```
void customer() {            // customer process
    down(&mutex);             // enter critical region
    if (waiting < CHAIRS){    // when seats available
        waiting++;           // one more waiting
        up(&customers);       // tell barber if first customer
        up(&mutex);           // release 'waiting'
        down(&barbers);       // sleep if no barber available
        get_haircut();        // get serviced
    } else up(&mutex);        // shop is full, leave
}
```